

SPECIALIZATION PROJECT IN MATHEMATICAL SCIENCES

Title

Author:
SINDRE GRØSTAD

Supervisors:
Professor KNUT-ANDREAS LIE



NORWEGIAN UNIVERSITY OF SCIENCE AND TECHNOLOGY
FACULTY OF INFORMATION TECHNOLOGY AND ELECTRICAL ENGINEERING
DEPARTMENT OF MATHEMATICAL SCIENCES
November 16, 2018

Preface

Abstract

Table of Contents

Preface	i
Abstract	iii
Table of Contents	v
1 Introduction	1
2 Theory	3
2.1 Automatic differentiation	3
2.1.1 Forward Automatic Differentiation	4
2.1.2 Dual Numbers	5
2.1.3 Backward Automatic Differentiation	6
2.1.4 Forward Automatic Differentiation with multiple parameters	7
2.2 Julia	7
3 Implementation	9
3.1 Implementation of Automatic Differentiation	9
3.2 Applications of Automatic Differentiation	9
4 Result and discussion	11
5 Conclusion and Future Work	13
Bibliography	14
Appendix	17

Chapter 1

Introduction

Intro

Theory

2.1 Automatic differentiation

Automatic differentiation (AD) is a method that makes the computer derive the derivatives with very little effort from the user. If you have not heard of AD before, the first thing that you might think of is algebraic or symbolic differentiation. In this type of differentiation the computer learns the basic rules from calculus like e.g.

$$\begin{aligned}\frac{d}{dx}x^n &= n \cdot x^{n-1} \\ \frac{d}{dx}\cos(x) &= -\sin(x) \\ \frac{d}{dx}\exp x &= \exp x\end{aligned}$$

etc. and the chain- and product rule

$$\begin{aligned}\frac{d}{dx}f(x) \cdot g(x) &= f'(x) \cdot g(x) + f(x) \cdot g'(x) \\ \frac{d}{dx}f(g(x)) &= g'(x) \cdot f'(g(x)).\end{aligned}$$

The computer will then use these rules on symbolic variables to obtain the derivative of any function given. This will give perfectly accurate derivatives, but the disadvantage with this approach is that it is computational demanding and as $f(x)$ becomes more complex, the calculations will become slow.

If AD is not symbolic differentiation you might think that it is finite differences, where you use the definition of the derivative

$$\frac{df}{dx} = \frac{f(x+h) - f(x)}{h}$$

with a small h to obtain the numerical approximation of the derivative of f . This approach is not optimal because, first of all, if you choose an h too small, you will get problems with rounding errors on your computer. This is because when h is small, you will subtract two very similar numbers, $f(x+h)$ and $f(x)$ and then divide by a small number h . This means that any small rounding errors in the subtraction will be hugely magnified by the division. Secondly, if you choose h too large your approximation of the derivative will not be accurate.

AD can be split into two different methods - forward AD and backward AD. Both methods are similar to symbolic differentiation in the way that we implement the differentiation rules, but they differ by instead of differentiating symbols and then inserting values for the symbols, we keep track of the function values and the corresponding values of the derivatives as we go. Both methods do this by separating each expression into a finite set of elementary operations.

2.1.1 Forward Automatic Differentiation

In forward AD, the function value is stored in a tuple $[\cdot, \cdot]$. In this way, we can continuously update both the function value and the derivative value for every operation we perform on the function.

As an example, consider the function $f = f(x)$ with its derivative f_x where x is a scalar variable. Then the AD-variable x is the pair $[x, 1]$ and for f we have $[f, f_x]$. In the pair $[x, 1]$, x is the numerical value of x and $1 = \frac{dx}{dx}$. Similar for $f(x)$ where f is the numerical value of $f(x)$ and f_x the numerical value of $f'(x)$. We then define the arithmetic operators for such pairs such that for functions f and g ,

$$\begin{aligned} [f, f_x] \pm [g, g_x] &= [f \pm g, f_x \pm g_x], \\ [f, f_x] \cdot [g, g_x] &= [f \cdot g, f_x \cdot g + f \cdot g_x], \\ \frac{[f, f_x]}{[g, g_x]} &= \left[\frac{f}{g}, \frac{f_x \cdot g - f \cdot g_x}{g^2} \right]. \end{aligned} \tag{2.1}$$

It is also necessary to define the chain rule such that for a function $h(x)$

$$h(f(x)) = h([f, f_x]) = [h(f), f_x \cdot h'(f)].$$

The only thing that remains to define are the rules concerning elementary functions like

$$\begin{aligned} \exp([f, f_x]) &= [\exp(f), \exp(f) \cdot f_x], \\ \log([f, f_x]) &= [\log(f), \frac{f_x}{f}], \\ \sin([f, f_x]) &= [\sin(f), \sin(f) \cdot f_x], \text{ etc.} \end{aligned} \tag{2.2}$$

When these arithmetic operators and the elementary function are implemented you are able to calculate any scalar function derivative without actually doing any form of differentiation yourselves. Let us look at an step by step example where

$$f(x) = x \cdot \exp(2x) \quad \text{for } x = 2. \tag{2.3}$$

Then the declaration of the AD-variable x gives $x = [2, 1]$. All scalars can be looked at as AD variables with derivative equal to 0 such that

$$\begin{aligned} 2x &= [2, 0] \cdot [2, 1] \\ &= [2 \cdot 2, 0 \cdot 1 + 2 \cdot 1] \\ &= [4, 2]. \end{aligned}$$

After this computation we get from the exponential

$$\begin{aligned}\exp(2x) &= \exp([4, 2]) \\ &= [\exp(4), \exp(4) \cdot 2],\end{aligned}$$

and lastly from the product rule we get the correct tuple for $f(x)$

$$\begin{aligned}x \cdot \exp(2x) &= [2, 1] \cdot [\exp(4), 2 \cdot \exp(4)] \\ &= [2 \cdot \exp(4), 1 \cdot \exp(4) + 2 \cdot 2 \exp(4)] \\ [f, f_x] &= [2 \cdot \exp(4), 5 \cdot \exp(4)].\end{aligned}$$

This result is equal

$$(f(x), f_x(x)) = (x \cdot \exp(2x), (1 + 2x) \exp(2x))$$

for $x = 2$.

2.1.2 Dual Numbers

One approach to implementing forward AD is by dual numbers. Similar to complex numbers dual numbers are defined as

$$a + b\epsilon. \quad (2.4)$$

Here a and b are scalars and corresponds to the function value and the derivative value. ϵ is like we have for complex numbers $i^2 = -1$, but the corresponding relation for dual numbers are $\epsilon^2 = 0$. The convenient part of implementing forward AD with dual numbers is that you get the differentiation rules for arithmetic operations for free. Consider the dual numbers x and y on the form of definition (2.4). Then we get for addition

$$\begin{aligned}x + y &= (a + b\epsilon) \cdot (c + d\epsilon) \\ &= a + c + (b + d)\epsilon,\end{aligned}$$

for multiplication

$$\begin{aligned}x \cdot y &= (a + b\epsilon) \cdot (c + d\epsilon) \\ &= ac + (ad + bc)\epsilon + bd\epsilon^2 \\ &= ac + (ad + bc)\epsilon,\end{aligned}$$

and for division

$$\begin{aligned}\frac{x}{y} &= \frac{a + b\epsilon}{c + d\epsilon} \\ &= \frac{a + b\epsilon}{c + d\epsilon} \cdot \frac{c - d\epsilon}{c - d\epsilon} \\ &= \frac{ac - (ad - bc)\epsilon - bd\epsilon^2}{c^2 - d\epsilon^2} \\ &= \frac{a}{c} + \frac{bc - ad}{c^2} \epsilon.\end{aligned}$$

This is very convenient, but how does dual numbers handle elementary functions like sin, exp, log etc? If we look at the Taylor expansion of a function $f(x)$ where x is a dual number, we get

$$\begin{aligned} f(x) = f(a + b\epsilon) &= f(a) + \frac{f'(a)}{1!}(b\epsilon) + \frac{f''(a)}{2!}(b\epsilon)^2 + \dots \\ &= f(a) + f'(a)b\epsilon. \end{aligned}$$

This means that to make dual numbers handle elementary functions, the first order Taylor expansion needs to be implemented. This equals the implementation of elementary differentiation rules described in equations (2.2).

The weakness of implementing AD with dual numbers is clear for functions with multiple variables. Let the function f be defined as $f(x, y, z) = x \cdot y + z$. Let us say we want to know the function value for $(x, y, z) = (2, 3, 4)$ together with all the derivatives of f . First we evaluate f with x as the only varying parameter and the rest as constants:

$$\begin{aligned} f(x, y, z) &= (2 + 1\epsilon) \cdot (3 + 0\epsilon) + (1 + 0\epsilon) \\ &= 7 + 3\epsilon. \end{aligned}$$

7 is now the function value of f , while 3 is the derivative value of f with respect to x , f_x . To obtain f_y and f_z we need two more function evaluations with respectively y and z as the varying parameters. This example illustrates the weakness of forward AD implemented with dual numbers - when the function evaluated have many input variables, we need equally many function evaluations to determine the jacobian of the function.

2.1.3 Backward Automatic Differentiation

The main disadvantage with forward AD is when there are many input variables and you want the derivative with respect to all variables. This is where Backward AD is a more efficient way of obtaining the derivatives. To explain backward AD it is easier to first consider the approach for forward AD, where the method also can be explained as an extensive use of the chain rule

$$\frac{\partial f}{\partial x} = \sum_i \left(\frac{\partial f}{\partial u_i} \cdot \frac{\partial u_i}{\partial x} \right). \quad (2.5)$$

Take $f(x) = x \cdot \exp(2x)$ like in the forward AD example (2.3). Then we split up the expression into functions w_1 , w_2 , w_3 and f such that

$$w_1 = x \quad w_2 = 2w_1 \quad w_3 = \exp(w_2)$$

Now the function can be described as $f = w_1 \cdot w_3$. If we want the derivative of f with respect to x we can obtain expressions for all w 's by using the chain rule (2.5)

$$\begin{aligned} \frac{\partial w_1}{\partial x} &= 1, \\ \frac{\partial w_2}{\partial x} &= 2 \cdot \frac{\partial w_1}{\partial x} = 2, \\ \frac{\partial w_3}{\partial x} &= \frac{\partial}{\partial w_2} \exp(w_2) \cdot \frac{\partial w_2}{\partial x} = 2 \exp(2x). \end{aligned}$$

Lastly by calculating the derivative of f with respect to x in the same matter yields

$$\begin{aligned}\frac{\partial f}{\partial x} &= \frac{\partial w_1}{\partial x} \cdot w_3 + w_1 \cdot \frac{\partial w_3}{\partial x} \\ &= \exp(2x) + x \cdot 2 \exp(2x) \\ &= (1 + 2x) \exp(2x).\end{aligned}$$

This shows that forward AD split up the operations and uses the chain rule with respect to the dependent variables. Backward AD also uses the chain rule, but in the opposite direction. It uses it with respect to the functions. The chain rule then has the form

$$\frac{\partial s}{\partial u} = \sum_i \left(\frac{\partial f_i}{\partial u} \cdot \frac{\partial s}{\partial f_i} \right).$$

TODO: Fix the notation and finish writing argument. ?See if it is possible to make changes to the forward AD example to make it more similar?

2.1.4 Forward Automatic Differentiation with multiple parameters

To handle the problem of many function evaluations in forward AD when having a function with many input parameters one can look at the method described by Knut-Andreas Lie in *User Guide for the MATLAB Reservoir Simulation Toolbox (MRST)*. Say we have three variables x , y and z . Then the corresponding AD-variables becomes

$$[x, (1, 0, 0)^T], \quad [y, (0, 1, 0)^T], \quad [z, (0, 0, 1)^T].$$

There are now not only one scalar as the derivative value, but the gradient to the corresponding variable. The operators defined in Equations 2.1 and the elementary functions in Equations 2.2 are still valid, but instead of scalar products there are now vector products. As an example let $f(x, y, z) = xyz$ and $x = 1$, $y = 2$ and $z = 3$, then

$$\begin{aligned}xyz &= [1, (1, 0, 0)^T] \cdot [2, (0, 1, 0)^T] \cdot [3, (0, 0, 1)^T] \\ &= [1 \cdot 2 \cdot 3, 2 \cdot 3 \cdot (1, 0, 0)^T + 1 \cdot 3 \cdot (0, 1, 0)^T + 1 \cdot 2 \cdot (0, 0, 1)^T] \\ [f, f_x] &= [6, (6, 3, 2)^T].\end{aligned}$$

This result is equal to the tuple

$$(f(x, y, z), \nabla f(x, y, z)) = (xyz, (yz, xz, xy)^T)$$

for the corresponding x , y and z values.

2.2 Julia

Er vel relevant å ha en seksjon om Julia, men i teori eller annet sted? I hvilken grad skal jeg skrive om at det er en implementasjon av noe som allerede er gjort i Matlab i MRST, men som nå implementeres i Julia? Hvor er det eventuelt relevant å skrive om det?

Implementation

3.1 Implementation of Automatic Differentiation

When it comes to implementing Automatic Differentiation(AD) there are two major concerns. First is that it must be easy and intuitive to use, the second is that it must be efficient code as it will be used in computational demanding calculations.

A convenient way to store the AD-variables in Julia is to make a struct that have two member variables, `val` and `jac`, that stores respectively the value and the corresponding Jacobian. The importance of the way you implement the AD operators can be expressed in a short example: Consider you have two variables x and y and you want to compute the function $f(x, y) = y + \exp(2xy)$. If the implementation is based on making new functions that take in AD-variables as input parameters, it will look something like this:

$$f = \text{ADplus}(y, \text{ADexp}(\text{ADtimes}(2, \text{ADtimes}(x, y)))).$$

This is clearly not a suitable way to implement AD and should be avoided. Instead of making new functions that takes in AD-variables as parameters one should overload the standard operators (+, -, *, /) and the elementary functions (exp, sin, log, etc.). In Julia this involves overloading the Base module such that when you write $x + y$ with x and y as AD-variables, Julia's Multiple Dispatch **ADD REF**, understand that it is your definition of the "+" operator that is meant to be used. This gives us the opportunity to only write $f = y + \exp(2xy)$ if we want to compute $f(x, y)$ for given x and y .

3.2 Applications of Automatic Differentiation

Next: starte med å skrive om newton-solveren jeg har laget som løser $f(x) = 0$.

Chapter 4

Result and discussion

Chapter 5

Conclusion and Future Work

Bibliography

User Guide for the MATLAB Reservoir Simulation Toolbox (MRST), 2018. An introduction to reservoir simulation using matlab/gnu octave. Accessed 25.10.2018.

URL <https://folk.ntnu.no/andreas/mrst/mrst-cam.pdf>

